

Recurrent Spiking Neural Networks

A quick overview of the e-prop online learning framework

Florent Pollet

Modelling in neuroscience
ENS Paris-Saclay

March 26, 2024



Introduction

Review of paper *A solution to the learning dilemma for recurrent networks of spiking neurons* Bellec et al., 2020, following papers Bellec et al., 2019 and Bellec et al., 2018.


- Artificial Neural Networks (ANN) consume more and more power
- Brain energy-efficiency is much higher than current electronic systems

⇒ Development of neuromorphic computing

- Spiking Neural Networks are a promising brain-inspired alternative to ANN
- However the most efficient methods to train them are not biologically plausible

⇒ How to design a biologically-plausible and efficient training strategy?

e-prop



"Neuromorphic computing"
by DALL-E 3

Outline

- 1 Recurrent Spiking Neural Networks
 - Neuron models
 - Coding
 - Architectures
 - Training
- 2 The e-prop learning framework
 - Biological considerations
 - e-prop definition and derivations
 - e-prop applications
- 3 Experiments with e-prop
 - Jax+Spyx implementation
 - e-prop on another dataset

Leaky Integrate-and-Fire (LIF) model - continuous

Analogy between ion and electron transfers, by Ramon Stein (1967):

- $V_m(t)$ membrane potential
- $I(t)$ input current
- $\tau_m = C_m R_m$ membrane time constant
- C_m membrane capacitor
- R_m membrane resistor
- Soft reset to V_{reset} after a spike emission (action potential) when V_{th} is crossed


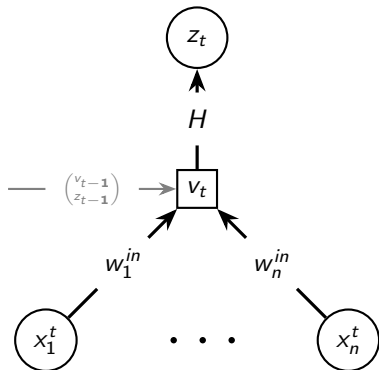


Figure: Equivalent circuit from Chen et al., 2021

$$\tau_m \frac{dV_m(t)}{dt} = -[V_m(t) - V_{rest}] + R_m I(t)$$

Leaky Integrate-and-Fire (LIF) model - discrete

- internal/hidden state v^t
- observable state $z^t \in \{0, 1\}$
- n inputs $(x_i^t)_i \in \{0, 1\}$
- w^{in} the input synapse weights
- δt the discrete time step size (1 ms for instance)
- $\alpha = \exp(-\delta t / \tau_m)$ the decay time constant
- H the Heaviside step function (non differentiable)



$$\begin{cases} v^{t+1} &= \alpha v^t + \sum_i w_i^{in} x_i^{t+1} - z^t v_{th} \\ z^{t+1} &= H(v^{t+1} - v_{th}) \end{cases}$$

Adaptive Leaky Integrate-and-Fire (ALIF) model

Introduced by Bellec et al., 2018, inspired by a biological phenomenon: spike-frequency adaptation (SFA)

Neuron with SFA \implies more expressivity for the neural network

- a_t the deviation in the baseline threshold voltage
- β scaling factor
- $\rho = \exp(-\delta t / \tau_a)$ the exponential decay parameter
- τ_a the adaptation time constant of the threshold decay

$$\begin{cases} v^{t+1} &= \alpha v^t + \sum_i w_i^{in} x_i^{t+1} - z^t v_{th} \\ a^{t+1} &= \rho a^t + z^t \\ z^{t+1} &= H(v^{t+1} - v_{th} - \beta a^t) \end{cases}$$

Refractory period can be enforced for both LIF and ALIF models by setting z_t to 0 for a few timesteps

LIF - ALIF comparison

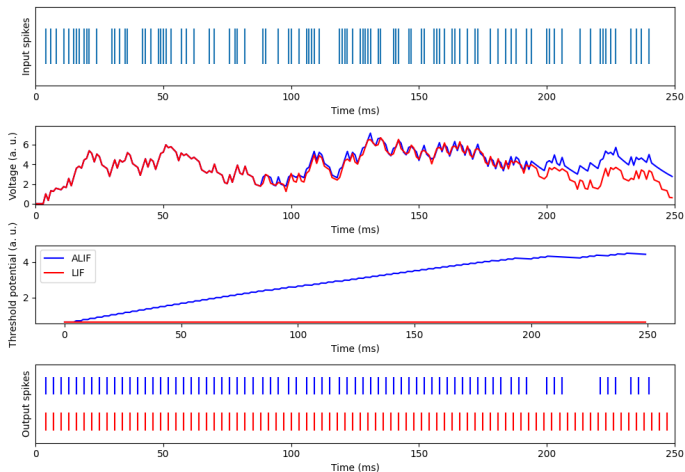


Figure: Comparison of LIF and ALIF neuron models with a random input spike train

Other famous neuron models

- Izhikevich: includes the refractory period in the model, but shown to perform badly with e-prop in van der Veen, 2022
- Hodgkin-Huxley: most biologically accurate, conductance instead of current input based, but more computationally-intensive
- Leaky output neurons

$$y^t = \kappa y^{t-1} + \sum_j w_j z_j^t + b \quad (1)$$

- Real-valued, not spiking
- Decay factor $\kappa = \exp(-\delta_t/\tau_{out})$ by analogy to the LIF neurons
- Bias b
- Code implementation: $y^t = \kappa y^{t-1} + (1 - \kappa) \sum_j w_j z_j^t$

Input data encoding

- Rate encoding:
 - Widespread in the community
 - Not the sparsest (not ideal for energy efficiency)
- Temporal encoding
 - Very sparse
 - Can resort to logscale to increase the range
- Other coding schemes
 - Delta modulation coding, inspired from neurophysiology
- Decoding can be done in a similar fashion

Input data encoding

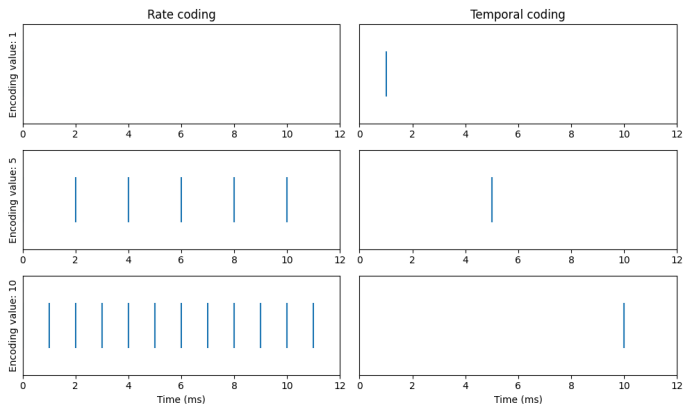


Figure: Rate and temporal coding comparison for simple integer values

SNN architectures

- Classical feed-forward architectures
 - Can replicate famous ANNs like ResNet, VGG
 - More common, mainly for Computer Vision or tasks related to Robotics
- Recurrent architectures (RSNN)
 - Closer to brain networks
 - The paper Bellec et al., 2018 introduces the Long short-term memory Spiking Neural Networks (LSNN), which is an RNN with a mix of LIF and ALIF neurons, thus showing SFA
- Transformers: early development Zhou et al., 2022

LSNN computational graph

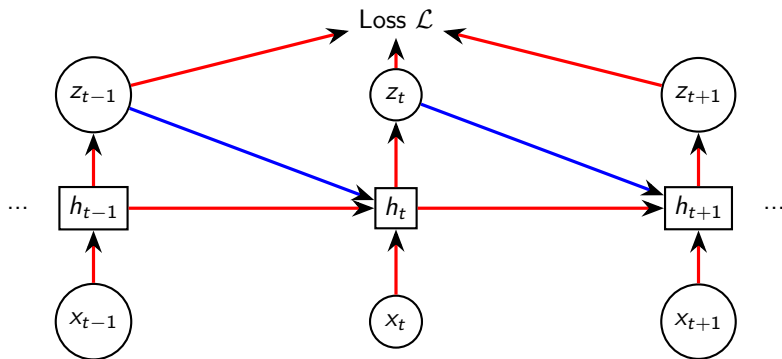


Figure: Example of a computational graph for a Spiking Neural Network. Inputs are written $(x_t)_t$, outputs $(z_t)_t$, and hidden states of the network $(h_t)_t$. Blue arrows are only present for Recurrent Neural Networks.

SNN training strategies

Main issue: activation functions are not differentiable

- ANN-to-SNN conversion: easier approach and best results for Deep SNN
- Spike-timing-dependent plasticity (STDP)
 - Most biologically-plausible method
 - Variation of Hebbian learning *Neurons who fire together, wire together*
 - Temporal order taken into account (pre- and post-synaptic)
 - Medium performance
- BackPropagation Through Time Werbos, 1990
 - Adapted from RNN/LSTM training in Deep Learning
 - Uses surrogate gradients like arctan, Gaussian functions or even Straight-Through Estimators
 - Good performance
- Real-Time Recurrent Learning
 - Online method
 - Bad complexity $O(n^4)$, with n the number of neurons

Biological considerations

In the brain:

- Eligibility traces: preceding activity leaves traces at molecular level (calcium ions or CAMKII enzymes)
- Learning signals: up down signals thanks to dopamine or acetylcholine, to inform the neurons of behavioral results



Figure: Credits <https://around.uoregon.edu/content/uo-neuroscientists-get-new-view-how-neurons-communicate>

e-prop main equations

We define the ideal learning signal for neuron j as L_j^t and the eligibility trace for synapse of indices ji as e_{ji}^t :

- $\frac{d}{d\cdot}$ indicates the total derivative
- $\frac{\partial}{\partial\cdot}$ indicates the partial derivative, which only makes sense between directly connected nodes of the computational graph
- $[\frac{d}{d\cdot}]_{local}$ consists in the total derivative in an altered computational graph, for which links from z_t to h_{t+1} are removed: see next slide

$$L_j^t = \frac{\partial \mathcal{L}}{\partial z_j^t} \quad (\approx \frac{d\mathcal{L}}{dz_j^t})$$

$$\begin{aligned} e_{ji}^t &= \left[\frac{dz_j^t}{dW_{ji}} \right]_{local} \\ &= \frac{\partial z_j^t}{\partial h_j^t} \cdot \sum_{t' \leq t} \frac{\partial h_j^t}{\partial h_j^{t-1}} \cdots \frac{\partial h_j^{t+1}}{\partial h_j^{t'}} \cdot \frac{\partial h_j^{t'}}{\partial W_{ji}} \end{aligned}$$

e-prop derivations - truncated graph

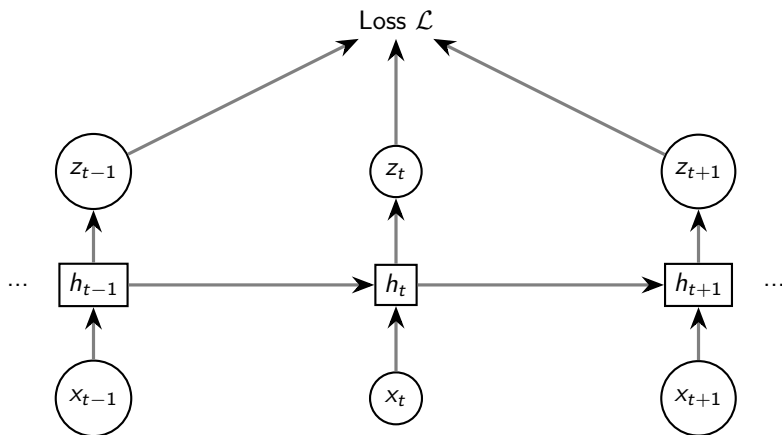


Figure: Truncated graph for local derivatives

e-prop derivations - chain rule

From the BPTT original paper Werbos, 1990, we have this factorization of the derivatives:

$$\frac{d\mathcal{L}}{dW_{ji}} = \sum_{t'} \frac{d\mathcal{L}}{dh_j^{t'}} \cdot \frac{\partial h_j^{t'}}{\partial W_{ji}}$$

We can expand this expressive using a recursive formula:

$$\frac{d\mathcal{L}}{dh_j^t} = \frac{d\mathcal{L}}{dz_j^t} \cdot \frac{\partial z_j^t}{\partial h_j^t} + \frac{d\mathcal{L}}{dh_j^{t+1}} \cdot \frac{\partial h_j^{t+1}}{\partial h_j^t}$$

See next slide for to visualize on the graph

e-prop derivations - chain rule

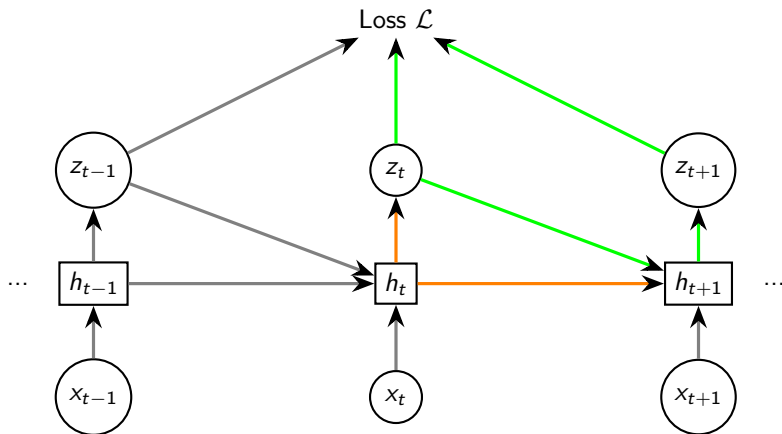


Figure: Chain rule for e-prop derivations $\frac{d\mathcal{L}}{dh_j^t}$

e-prop derivations

We can therefore write the full recursion and reorder the sum, to get the e-prop main equation:

$$\begin{aligned}
 \frac{d\mathcal{L}}{dW_{ji}} &= \sum_{1 \leq t' \leq t_m} \left(\sum_{t' \leq t \leq t_m} \frac{d\mathcal{L}}{dz_j^t} \cdot \frac{\partial z_j^t}{\partial h_j^t} \cdot \left[\prod_{t'+1 \leq t'' \leq t} \frac{\partial h_j^{t''}}{\partial h_j^{t''-1}} \right] \right) \cdot \frac{\partial h_j^{t'}}{\partial W_{ji}} \\
 &= \sum_{1 \leq t \leq t_m} \frac{d\mathcal{L}}{dz_j^t} \cdot \frac{\partial z_j^t}{\partial h_j^t} \cdot \left(\sum_{1 \leq t' \leq t} \left[\prod_{t'+1 \leq t'' \leq t} \frac{\partial h_j^{t''}}{\partial h_j^{t''-1}} \right] \cdot \frac{\partial h_j^{t'}}{\partial W_{ji}} \right) \\
 &= \sum_t \frac{d\mathcal{L}}{dz_j^t} e_{ji}^t
 \end{aligned}$$

To enforce an online setting, one approximate $\frac{d\mathcal{L}}{dz_j^t}$ by its partial derivative (see next slide).

e-prop derivations - learning signal approximation




Figure: Learning signal approximation for e-prop derivations (orange: partial derivative, orange+green: total derivative)

Broadcast weights

Supervised learning here (for Reinforcement Learning: reward e-prop):

$$\mathcal{L}_{MSE} = \frac{1}{2} \sum_{t,k} (y_k^t - y_k^{*,t})^2 \quad \mathcal{L}_{CE} = - \sum_{t,k} \pi_k^{*,t} \log \pi_k^t$$

Family of learning signals with broadcast weights (B_{jk}):

$$\tilde{l}_j^t = \sum_k B_{jk} \sum_{t' \geq t} (\zeta_k^{t'} - \zeta_k^{*,t'}) \kappa^{t'-t}$$

Hence, the e-prop gradient:

$$\begin{aligned} \frac{d\mathcal{L}}{dW_{ji}} &\approx \sum_{k,t} B_{jk} \sum_{t' \geq t} (\zeta_k^{t'} - \zeta_k^{*,t'}) \kappa^{t'-t} e_{ji}^t \\ &\approx \sum_{k,t} B_{jk} (\zeta_k^t - \zeta_k^{*,t}) \sum_{t' \leq t} \kappa^{t-t'} e_{ji}^{t'} \end{aligned}$$

Temporal low-pass filter of the eligibility traces

e-prop variants

How to choose the broadcast weights?

- symmetric e-prop (ideal learning signal, best approximation): $B_{jk} = W_{kj}^{out}$
- random e-prop: constant random weights bad performance, except for multi-layer
- adaptive e-prop: random weights updated with the same rule as the other weights

Detailed calculations for each method can be found in Bellec et al., 2020.

Learning phoneme recognition

- TIMIT dataset from Garofolo, John S. et al., 1993
- Preprocessing to extract Mel coefficients, as inputs
- No encoding, continuous input in $[0, 1]$




Figure: TIMIT tasks




Figure: TIMIT results

Difficult temporal credit assignment

- Goal: choose the side with most cues
- Interest: prove the learning of long-time dependencies


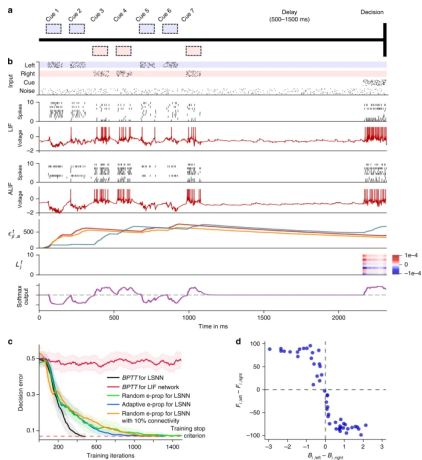


Figure: Architecture comparison

Figure: Task results

symmetric e-prop through auto-differentiation

How to leverage auto-differentiation from modern machine learning frameworks to compute e-prop?

By stopping the gradients from z_t to h_{t+1} , we obtain the truncated computational graph shown in the previous slide, that leads to the following simplifications:

$$\begin{aligned}L_j^t &= \frac{\partial \mathcal{L}}{\partial z_j^t} = \frac{d\mathcal{L}}{dz_j^t} \\e_{ji}^t &= \frac{dz_j^t}{dW_{ji}} \\ \frac{d\mathcal{L}}{dW_{ji}} &= \sum_t L_j^t e_{ji}^t\end{aligned}$$

\implies We can use auto-differentiation to compute symmetric e-prop.

Why Jax instead of PyTorch? It can be XLA JIT-compiled and the code can fully run on GPU with optimized kernels.

Jax+Spyx implementation: gradient computations



Figure: Jax logo - Bradbury et al., 2018




Figure: Spyx logo - Heckel and Nowotny, 2024

	EH vs EA	EH vs BPTT autodiff
5 timesteps	e-08	0.1
25 timesteps	e-07	0.26
100 timesteps	e-07	0.57

Table: Norm of the differences between different gradients

EH = e-prop hardcoded

EA = e-prop autodiff

Spiking Heidelberg Digits (SHD) dataset

- Benchmark dataset for spiking neural networks
- Audio digit classification in German and English
- 1000 samples
- Delta modulation encoding

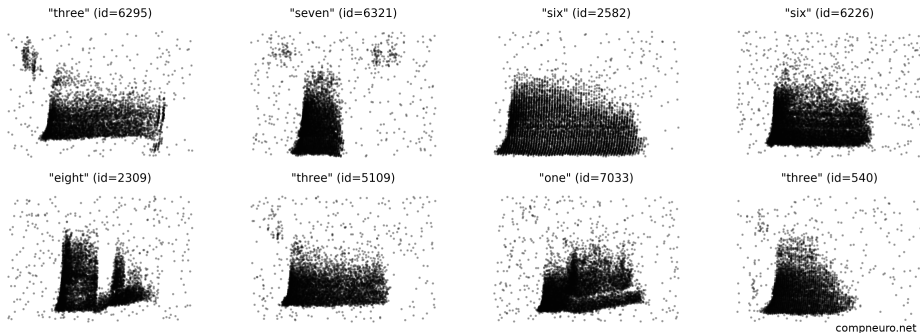


Figure: SHD dataset (x-axis: time, y-axis: channel)

LSNN with symmetric e-prop test on SHD dataset

- Network parameters: batch size of 256, 400 ALIF neurons, 60 iterations and 128 timesteps
- Training time: roughly 2 min (this is fast!)
- Benchmark test accuracy: 0.76

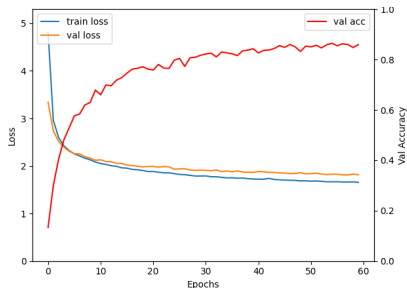


Figure: e-prop - train loss: 1.66, test accuracy: 0.75

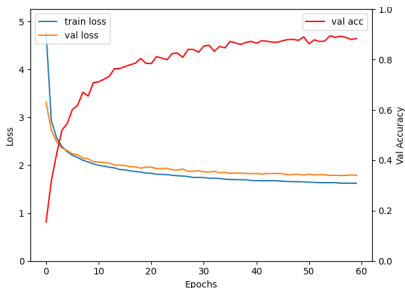


Figure: BPTT - train loss: 1.62, test accuracy: 0.76

LSNN with symmetric e-prop test on SHD dataset

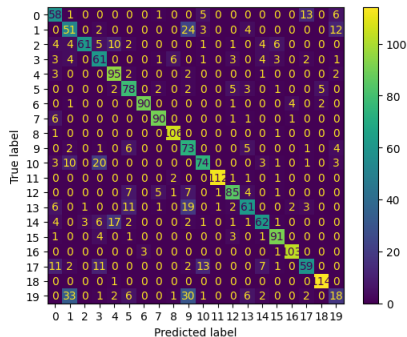


Figure: e-prop - train loss: 1.66, test accuracy: 0.75

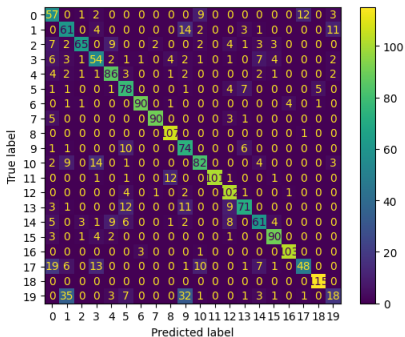


Figure: BPTT - train loss: 1.62, test accuracy: 0.76

Limitations and further work






- Connect this work to STDP, another biological learning mechanism, which is not discussed in the paper
- Compare e-prop to Real-Time Recurrent Learning (RTRL) and other variations, such as Online Training Through Time (OTTT) and Online SpatioTemporal Learning (OSTL) Summe, Schaefer, and Joshi, 2023 approximations
- e-prop may be more interesting for Reinforcement Learning (RL), with reward e-prop, because supervised learning is dominated by ANN
- Think about good heuristics for SNN design because of a very large hyperparameter search space compared to ANN
- Another interesting dataset to test Brain-To-Text, with real human spiking data <https://eval.ai/web/challenges/challenge-page/2099/>

Thanks! Any questions?







Code: <https://github.com/florian6973/btt-spyx>

Could the report be uploaded to arXiv?

References I

-  Bellec, Guillaume et al. (2018). *Long short-term memory and learning-to-learn in networks of spiking neurons*. [arXiv: 1803.09574 \[cs.NE\]](#).
-  Bellec, Guillaume et al. (2019). *Biologically inspired alternatives to backpropagation through time for learning in recurrent neural nets*. [arXiv: 1901.09049 \[cs.NE\]](#).
-  Bellec, Guillaume et al. (July 2020). "A solution to the learning dilemma for recurrent networks of spiking neurons". In: *Nature Communications* 11.1. ISSN: 2041-1723. DOI: 10.1038/s41467-020-17236-y. URL: <http://dx.doi.org/10.1038/s41467-020-17236-y>.
-  Bradbury, James et al. (2018). *JAX: composable transformations of Python+NumPy programs*. Version 0.3.13. URL: <http://github.com/google/jax>.
-  Chen, Jiankun et al. (2021). *SAR Image Classification Based on Spiking Neural Network through Spike-Time Dependent Plasticity and Gradient Descent*. [arXiv: 2106.08005 \[cs.CV\]](#).

References II

-  Garofolo, John S. et al. (1993). *TIMIT Acoustic-Phonetic Continuous Speech Corpus*. DOI: 10.35111/17GK-BN40. URL: <https://catalog.ldc.upenn.edu/LDC93S1>.
-  Heckel, Kade M. and Thomas Nowotny (2024). *Spyx: A Library for Just-In-Time Compiled Optimization of Spiking Neural Networks*. arXiv: 2402.18994 [cs.NE].
-  Summe, Thomas, Clemens JS Schaefer, and Siddharth Joshi (2023). *Estimating Post-Synaptic Effects for Online Training of Feed-Forward SNNs*. arXiv: 2311.16151 [cs.NE].
-  Van der Veen, Werner (2022). *Including STDP to eligibility propagation in multi-layer recurrent spiking neural networks*. arXiv: 2201.07602 [cs.NE].
-  Werbos, P.J. (1990). "Backpropagation through time: what it does and how to do it". In: *Proceedings of the IEEE* 78.10, pp. 1550–1560. DOI: 10.1109/5.58337.
-  Zhou, Zhaokun et al. (2022). *Spikformer: When Spiking Neural Network Meets Transformer*. arXiv: 2209.15425 [cs.NE].